

Google Vizier: A Service for Black-Box Optimization

Daniel Golovin
Google Research
dgg@google.com

Benjamin Solnik
Google Research
bsolnik@google.com

Subhodeep Moitra
Google Research
smoitra@google.com

Greg Kochanski
Google Research
gpk@google.com

John Karro
Google Research
karro@google.com

D. Sculley
Google Research
dsculley@google.com

ABSTRACT

Any sufficiently complex system acts as a black box when it becomes easier to experiment with than to understand. Hence, black-box optimization has become increasingly important as systems have become more complex. In this paper we describe *Google Vizier*, a Google-internal service for performing black-box optimization that has become the de facto parameter tuning engine at Google. **Google Vizier is used to optimize many of our machine learning models and other systems, and also provides core capabilities to Google's Cloud Machine Learning HyperTune subsystem.** We discuss our requirements, infrastructure design, underlying algorithms, and advanced features such as transfer learning and automated early stopping that the service provides.

KEYWORDS

Black-Box Optimization, Bayesian Optimization, Gaussian Processes, Hyperparameters, Transfer Learning, Automated Stopping

1 INTRODUCTION

Black-box optimization is the task of optimizing an objective function $f : X \rightarrow \mathbb{R}$ with a limited budget for evaluations. The adjective “black-box” means that while we can evaluate $f(x)$ for any $x \in X$, we have no access to any other information about f , such as gradients or the Hessian. When function evaluations are expensive, it makes sense to carefully and adaptively select values to evaluate; the overall goal is for the system to generate a sequence of x_t that approaches the global optimum as rapidly as possible.

Black box optimization algorithms can be used to find the best operating parameters for any system whose performance can be measured as a function of adjustable parameters. It has many important applications, such as automated tuning of the hyperparameters of machine learning systems (e.g., learning rates, or the number of hidden layers in a deep neural network), optimization of the user interfaces of web services (e.g. optimizing colors and fonts

to maximize reading speed), and optimization of physical systems (e.g., optimizing airfoils in simulation).

In this paper we discuss a state-of-the-art system for black-box optimization developed within Google, called *Google Vizier*, named after a high official who offers advice to rulers. It is a service for black-box optimization that supports several advanced algorithms. The system has a convenient **Remote Procedure Call (RPC) interface**, along with a dashboard and analysis tools. Google Vizier is a research project, parts of which supply core capabilities to our Cloud Machine Learning *HyperTune*¹ subsystem. We discuss the architecture of the system, design choices, and some of the algorithms used.

1.1 Related Work

Black-box optimization makes minimal assumptions about the problem under consideration, and thus is broadly applicable across many domains and has been studied in multiple scholarly fields under names including Bayesian Optimization [2, 25, 26], Derivative-free optimization [7, 24], Sequential Experimental Design [5], and assorted variants of the multiarmed bandit problem [13, 20, 29].

Several classes of algorithms have been proposed for the problem. The simplest of these are non-adaptive procedures such as RANDOM SEARCH, which selects x_t uniformly at random from X at each time step t independent of the previous points selected, $\{x_\tau : 1 \leq \tau < t\}$, and GRID SEARCH, which selects along a grid (i.e., the Cartesian product of finite sets of feasible values for each parameter). Classic algorithms such as SIMULATED ANNEALING and assorted genetic algorithms have also been investigated, e.g., Covariance Matrix Adaptation [16].

Another class of algorithms performs a local search by selecting points that maintain a search pattern, such as a simplex in the case of the classic NELDER–MEAD algorithm [22]. More modern variants of these algorithms maintain simple models of the objective f within a subset of the feasible regions (called the *trust region*), and select a point x_t to improve the model within the trust region [7].

More recently, some researchers have combined powerful techniques for modeling the objective f over the entire feasible region, using ideas developed for multiarmed bandit problems for managing explore / exploit trade-offs. These approaches are fundamentally Bayesian in nature, hence this literature goes under the name *Bayesian Optimization*. Typically, the model for f is a Gaussian



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivs International 4.0 License.

KDD '17, August 13–17, 2017, Halifax, NS, Canada
© 2017 Copyright held by the owner/author(s).
ACM ISBN 978-1-4503-4887-4/17/08.
<https://doi.org/10.1145/3097983.3098043>

¹<https://cloud.google.com/ml/>

process (as in [26, 29]), a deep neural network (as in [27, 31]), or a regression forest (as in [2, 19]).

Many of these algorithms have open-source implementations available. Within the machine learning community, examples include, e.g., HyperOpt², MOE³, Spearmint⁴, and AutoWeka⁵, among many others. In contrast to such software packages, which require practitioners to set them up and run them locally, we opted to develop a managed service for black-box optimization, which is more convenient for users but involves additional design considerations.

1.2 Definitions

Throughout the paper, we use the following terms to describe the semantics of the system:

A *Trial* is a list of parameter values, x , that will lead to a single evaluation of $f(x)$. A trial can be “Completed”, which means that it has been evaluated and the objective value $f(x)$ has been assigned to it, otherwise it is “Pending”.

A *Study* represents a single optimization run over a feasible space. Each Study contains a configuration describing the feasible space, as well as a set of Trials. It is assumed that $f(x)$ does not change in the course of a Study.

A *Worker* refers to a process responsible for evaluating a Pending Trial and calculating its objective value.

2 SYSTEM OVERVIEW

This section explores the design considerations involved in implementing black-box optimization as a service.

2.1 Design Goals and Constraints

Vizier’s design satisfies the following desiderata:

- Ease of use. Minimal user configuration and setup.
- Hosts state-of-the-art black-box optimization algorithms.
- High availability
- Scalable to millions of trials per study, thousands of parallel trial evaluations per study, and billions of studies.
- Easy to experiment with new algorithms.
- Easy to change out algorithms deployed in production.

For ease of use, we implemented Vizier as a managed service that stores the state of each optimization. This approach drastically reduces the effort a new user needs to get up and running; and a managed service with a well-documented and stable RPC API allows us to upgrade the service without user effort. We provide a default configuration for our managed service that is good enough to ensure that most users need never concern themselves with the underlying optimization algorithms.

The default option allows the service to dynamically select a recommended black-box algorithm along with low-level settings based on the study configuration. We choose to make our algorithms stateless, so that we can seamlessly switch algorithms during a

study, dynamically choosing the algorithm that is likely to perform better for a particular trial of a given study. For example, Gaussian Process Bandits [26, 29] provide excellent result quality, but naive implementations scale as $O(n^3)$ with the number of training points. Thus, once we’ve collected a large number of completed Trials, we may want to switch to using a more scalable algorithm.

At the same time, we want to allow ourselves (and advanced users) the freedom to experiment with new algorithms or special-case modifications of the supported algorithms in a manner that is safe, easy, and fast. Hence, we’ve built Google Vizier as a modular system consisting of four cooperating processes (see Figure 1) that update the state of Studies in the central database. The processes themselves are modular with several clean abstraction layers that allow us to experiment with and apply different algorithms easily.

Finally we want to allow multiple trials to be evaluated in parallel, and allow for the possibility that evaluating the objective function for each trial could itself be a distributed process. To this end we define Workers, responsible for evaluating suggestions, and identify each worker by a name (a `worker_handle`) that persists across process preemptions or crashes.

2.2 Basic User Workflow

To use Vizier, a developer may use one of our client libraries (currently implemented in C++, Python, Golang), which will generate service requests encoded as protocol buffers [15]. The basic workflow is extremely simple. Users specify a *study configuration* which includes:

- Identifying characteristics of the study (e.g. name, owner, permissions).
- The set of parameters along with feasible sets for each (c.f., Section 2.3.1 for details); Vizier does constrained optimization over the feasible set.

Given this configuration, basic use of the service (with each trial being evaluated by a single process) can be implemented as follows:

```
# Register this client with the Study, creating it if
# necessary.
client.LoadStudy(study_config, worker_handle)
while (not client.StudyIsDone()):
    # Obtain a trial to evaluate.
    trial = client.GetSuggestion()
    # Evaluate the objective function at the trial parameters.
    metrics = RunTrial(trial)
    # Report back the results.
    client.CompleteTrial(trial, metrics)
```

Here `RunTrial` is the problem-specific evaluation of the objective function f . Multiple named metrics may be reported back to Vizier, however one must be distinguished as the objective value $f(x)$ for trial x . Note that multiple processes working on a study should share the same `worker_handle` if and only if they are collaboratively evaluating the same trial. All processes registered with a given study with the same `worker_handle` are guaranteed to receive the same trial upon request, which enables distributed trial evaluation.

²<https://github.com/jaberg/hyperopt>

³<https://github.com/Yelp/MOE>

⁴<https://github.com/HIPS/Spearmint>

⁵<https://github.com/automl/autoweka>

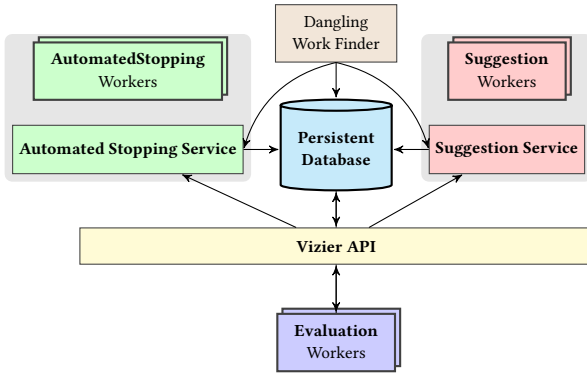


Figure 1: Architecture of Vizier service: Main components are (1) Dangling work finder (restarts work lost to preemptions) (2) Persistent Database holding the current state of all Studies (3) Suggestion Service (creates new Trials), (4) Early Stopping Service (helps terminate a Trial early) (5) Vizier API (JSON, validation, multiplexing) (6) Evaluation workers (provided and owned by the user).

2.3 Interfaces

2.3.1 Configuring a Study. To configure a study, the user provides a study name, owner, optional access permissions, an optimization goal from {MAXIMIZE, MINIMIZE}, and specifies the feasible region X via a set of `ParameterConfigs`, each of which declares a parameter name along with its values. We support the following parameter types:

- **DOUBLE:** The feasible region is a closed interval $[a, b]$ for some real values $a \leq b$.
- **INTEGER:** The feasible region has the form $[a, b] \cap \mathbb{Z}$ for some integers $a \leq b$.
- **DISCRETE:** The feasible region is an explicitly specified, ordered set of real numbers.
- **CATEGORICAL:** The feasible region is an explicitly specified, unordered set of strings.

Users may also suggest recommended scaling, e.g., logarithmic scaling for parameters for which the objective may depend only on the order of magnitude of a parameter value.

2.3.2 API Definition. Workers and end users can make calls to the Vizier Service using either a REST API or using Google’s internal RPC protocol [15]. The most important service calls are:

- **CreateStudy:** Given a Study configuration, this creates an optimization Study and returns a globally unique identifier (“guid”) which is then used for all future service calls. If a Study with a matching name exists, the guid for that Study is returned. This allows parallel workers to call this method and all register with the same Study.
- **SuggestTrials:** This method takes a “worker handle” as input, and immediately returns a globally unique handle for a “long-running operation” that represents the work of generating Trial suggestions. The user can then poll the API periodically to check the status of the operation. Once the operation is completed, it will contain the suggested Trials. This design

ensures that all service calls are made with low latency, while allowing for the fact that the generation of Trials can take longer.

- **AddMeasurementToTrial:** This method allows clients to provide intermediate metrics during the evaluation of a Trial. These metrics are then used by the Automated Stopping rules to determine which Trials should be stopped early.
- **CompleteTrial:** This method changes a Trial’s status to “Completed”, and provides a final objective value that is then used to inform the suggestions provided by future calls to `SuggestTrials`.
- **ShouldTrialStop:** This method returns a globally unique handle for a long-running operation that represents the work of determining whether a Pending Trial should be stopped.

2.4 Infrastructure

2.4.1 Parallel Processing of Suggestion Work. As the de facto parameter tuning engine of Google, Vizier is constantly working on generating suggestions for a large number of Studies concurrently. As such, a single machine would be insufficient for handling the workload. Our Suggestion Service is therefore partitioned across several Google datacenters, with a number of machines being used in each one. Each instance of the Suggestion Service potentially can generate suggestions for several Studies in parallel, giving us a massively scalable suggestion infrastructure. Google’s load balancing infrastructure is then used to allow clients to make calls to a unified endpoint, without needing to know which instance is doing the work.

When a request is received by a Suggestion Service instance to generate suggestions, the instance first places a distributed lock on the Study. This lock is acquired for a fixed period of time, and is periodically extended by a separate thread running on the instance. In other words, the lock will be held until either the instance fails, or it decides it’s done working on the Study. If the instance fails (due to e.g. hardware failure, job preemption, etc), the lock soon expires, making it eligible to be picked up by a separate process (called the “DanglingWorkFinder”) which then reassigns the Study to a different Suggestion Service instance.

One consideration in maintaining a production system is that bugs are inevitably introduced as our code matures. Occasionally, a new algorithmic change, however well tested, will lead to instances of the Suggestion Service failing for particular Studies. If a Study is picked up by the DanglingWorkFinder too many times, it will temporarily halt the Study and alert us. This prevents subtle bugs that only affect a few Studies from causing crash loops that affect the overall stability of the system.

2.5 The Algorithm Playground

Vizier’s *algorithm playground* provides a mechanism for advanced users to easily, quickly, and safely replace Vizier’s core optimization algorithms with arbitrary algorithms.

The playground serves a dual purpose; it allows rapid prototyping of new algorithms, and it allows power-users to easily customize Vizier with advanced or exotic capabilities that are particular to




Figure 2: Architecture of Playground mode: Main components are (1) The Vizier API takes service requests. (2) The Custom Policy implements the Abstract Policy and generates suggested Trials. (3) The Playground Binary drives the custom policy based on demand reported by the Vizier API. (4) The Evaluation Workers behave as normal, i.e., they request and evaluate Trials.

their use-case. In all cases, users of the playground benefit from all of Vizier’s infrastructure aside from the core algorithms, such as access to a persistent database of Trials, the dashboard, and visualizations.

At the core of the playground is the ability to inject Trials into a Study. Vizier allows the user or other authorized processes to request one or more particular Trials be evaluated. In Playground mode, Vizier does not suggest Trials for evaluation, but relies on an external binary to generate Trials, which are then pushed to the service for later distribution to the workers.

More specifically, the architecture of the Playground involves the following key components: (1) *Abstract Policy* (2) *Playground Binary*, (3) *Vizier Service* and (4) *Evaluation Workers*. See Figure 2 for an illustration.

The *Abstract Policy* contains two abstract methods:

- (1) `GetNewSuggestions(trials, num_suggestions)`
- (2) `GetEarlyStoppingTrials(trials)`

which should be implemented by the user’s custom policy. Both these methods are passed the full state of all Trials in the Study, so algorithms may be implemented in a stateless fashion if desired. `GetNewSuggestions` is expected to generate `num_suggestions` new trials, while the `GetEarlyStoppingTrials` method is expected to return a list of Pending Trials that should be stopped early. The custom policy is registered with the *Playground Binary* which periodically polls the *Vizier Service*. The *Evaluation Workers* maintain the service abstraction and are unaware of the existence of the Playground.

2.6 Benchmarking Suite

Vizier has an integrated framework that allows us to efficiently benchmark our algorithms on a variety of objective functions. Many of the objective functions come from the Black-Box Optimization Benchmarking Workshop [10], but the framework allows for any function to be modeled by implementing an abstract *Experimenter* class, which has a virtual method responsible for calculating the objective value for a given Trial, and a second virtual method that returns the optimal solution for that benchmark.



Figure 3: A section of the dashboard for tracking the progress of Trials and the corresponding objective function values. Note also, the presence of actions buttons such as *Get Suggestions* for manually requesting suggestions.

Users configure a set of benchmark runs by providing a set of algorithm configurations and a set of objective functions. The benchmarking suite will optimize each function with each algorithm k times (where k is configurable), producing a series of performance-over-time metrics which are then formatted after execution. The individual runs are distributed over multiple threads and multiple machines, so it is easy to have thousands of benchmark runs executed in parallel.

2.7 Dashboard and Visualizations

Vizier has a web dashboard which is used for both monitoring and changing the state of Vizier studies. The dashboard is fully featured and implements the full functionality of the Vizier API. The dashboard is commonly used for: (1) Tracking the progress of a study; (2) Interactive visualizations; (3) Creating, updating and deleting a study; (4) Requesting new suggestions, early stopping, activating/deactivating a study. See Figure 3 for a section of the dashboard. In addition to monitoring and visualizations, the dashboard contains action buttons such as *Get Suggestions*.

The dashboard uses a translation layer which converts between JSON and protocol buffers [15] when talking with backend servers. The dashboard is built with Polymer [14] an open source web framework supported by Google and uses material design principles. It contains interactive visualizations for analyzing the parameters in your study. In particular, we use the parallel coordinates visualization [18] which has the benefit of scaling to high dimensional spaces (~15 dimensions) and works with both numerical and categorical parameters. See Figure 4 for an example. Each vertical axis is a dimension corresponding to a parameter, whereas each horizontal line is an individual trial. The point at which the horizontal line intersects the vertical axis gives the value of the parameter in that dimension. This can be used for examining how the dimensions co-vary with each other and also against the objective function value (left most axis). The visualizations are built using *d3.js* [4].




Figure 4: The Parallel Coordinates visualization [18] is used for examining results from different Vizier runs. It has the benefit of scaling to high dimensional spaces (~ 15 dimensions) and works with both numerical and categorical parameters. Additionally, it is interactive and allows various modes of slicing and dicing data.

3 THE VIZIER ALGORITHMS

Vizier’s modular design allows us to easily support **multiple algorithms**. For studies with under a thousand trials, Vizier defaults to using Batched Gaussian Process Bandits [8]. We use a Matérn kernel with automatic relevance determination (see e.g. section 5.1 of Rasmussen and Williams [23] for a discussion) and the *expected improvement* acquisition function [21]. We search for and find local maxima of the acquisition function with a proprietary gradient-free hill climbing algorithm, with random starting points.

We implement discrete parameters by embedding them in \mathbb{R} . Categorical parameters with k feasible values are represented via one-hot encoding, i.e., embedded in $[0, 1]^k$. In both cases, the Gaussian Process regressor gives us a continuous and differentiable function upon which we can walk uphill, then when the walk has converged, round to the nearest feasible point.

While some authors recommend using Bayesian deep learning models in lieu of Gaussian processes for scalability [27, 31], in our experience they are too sensitive to their own hyperparameters and do not reliably perform well. Other researchers have recognized this problem as well, and are working to address it [28].

For studies with tens of thousands of trials or more, other algorithms may be used. Though RANDOMSEARCH and GRIDSEARCH are supported as first-class choices and may be used in this regime, and many other published algorithms are supported through the algorithm playground, we currently recommend a proprietary local-search algorithm under these conditions.

For all of these algorithms we support data normalization, which maps numeric parameter values into $[0, 1]$ and objective values onto $[-0.5, 0.5]$. Depending on the problem, a one-to-one nonlinear mapping may be used for some of the parameters, and is typically used for the objective. Data normalization is handled before trials are presented to the trial suggestion algorithms, and its suggestions are transparently mapped back to the user-specified scaling.

3.1 Automated Early Stopping

In some important applications of black-box optimization, information related to the performance of a trial may become available during trial evaluation. Perhaps the best example of such a *performance curve* occurs when tuning machine learning hyperparameters for models trained progressively (e.g., via some version of stochastic gradient descent). In this case, the model typically becomes more accurate as it trains on more data, and the accuracy of the model is available at the end of each training epoch. Using these accuracy vs. training step curves, it is often possible to determine that a trial’s parameter settings are unpromising well before evaluation is finished. In this case we can terminate trial evaluation early, freeing those evaluation resources for more promising trial parameters. When done algorithmically, this is referred to as *automated early stopping*.

Vizier supports automated early stopping via an API call to a `ShouldTrialStop` method. Analogously to the Suggestion Service, there is an Automated Stopping Service that accepts requests from the Vizier API to analyze a study and determine the set of trials that should be stopped, according to the configured early stopping algorithm. As with suggestion algorithms, several automated early stopping algorithms are supported, and rapid prototyping can be done via the algorithm playground.

3.2 Automated Stopping Algorithms

Vizier supports the following automated stopping algorithms. These are meant to work in a stateless fashion i.e. they are given the full state of all trials in the Vizier study when determining which trials should stop.

3.2.1 Performance Curve Stopping Rule. This stopping rule performs regression on the performance curves to make a prediction of the final objective value of a Trial given a set of Trials that are already Completed, and a partial performance curve (i.e., a set of measurements taken during Trial evaluation). Given this prediction, if the probability of exceeding the optimal value found thus far is sufficiently low, early stopping is requested for the Trial.

While prior work on automated early stopping used Bayesian parametric regression [9, 30], we opted for a Bayesian non-parametric regression, specifically a Gaussian process model with a carefully designed kernel that measures similarity between performance curves. Our motivation in this was to be robust to many kinds of performance curves, including those coming from applications other than tuning machine learning hyperparameters in which the performance curves may have very different semantics. Notably, this stopping rule still works well even when the performance curve is not measuring the same quantity as the objective value, but is merely predictive of it.

3.2.2 Median Stopping Rule. The median stopping rule stops a pending trial x_t at step s if the trial’s best objective value by step s is strictly worse than the median value of the running averages $\hat{o}_{1:s}^T$ of all completed trials’ objectives x_τ reported up to step s . Here, we calculate the running average of a trial x_τ up to step s

as $\hat{o}_{1:s}^\tau = \frac{1}{s} \sum_{i=1}^s o_i^\tau$, where o_i^τ is the objective value of x_τ at step i . As with the performance curve stopping rule, the median stopping rule does not depend on a parametric model, and is applicable to a wide range of performance curves. In fact, the median stopping rule is model-free, and is more reminiscent of a bandit-based approach such as HYPERBAND [20].

3.3 Transfer learning

When doing black-box optimization, users often run studies that are similar to studies they have run before, and we can use this fact to minimize repeated work. **Vizier supports a form of Transfer Learning which leverages data from prior studies to guide and accelerate the current study. For instance, one might tune the learning rate and regularization of a machine learning system, then use that Study as a prior to tune the same ML system on a different data set.**

Vizier's current approach to transfer learning is relatively simple, yet robust to changes in objective across studies. We designed our transfer learning approach with these goals in mind:

- (1) Scale well to situations where there are many prior studies.
- (2) Accelerate studies (i.e., achieve better results with fewer trials) when the priors are good, particularly in cases where the location of the optimum, x^* , doesn't change much.
- (3) Be robust against poorly chosen prior studies (i.e., a bad prior should give only a modest deceleration).
- (4) Share information even when there is no formal relationship between the prior and current Studies.

In previous work on transfer learning in the context of hyperparameter optimization, Bardenet et al. [1] discuss the difficulty in transferring knowledge across different datasets especially when the observed metrics and the sampling of the datasets are different. They use a ranking approach for constructing a surrogate model for the response surface. This approach suffers from the computational overhead of running a ranking algorithm. Yogatama and Mann [32] propose a more efficient approach, which scales as $\Theta(kn + n^3)$ for k studies of n trials each, where the cubic term comes from using a Gaussian process in their acquisition function.

Vizier typically uses Gaussian Process regressors, so one natural approach to implementing transfer learning might be to build a larger Gaussian Process regressor that is trained on both the prior(s) and the current Study. However that approach fails to satisfy design goal 1: for k studies with n trials each it would require $\Omega(k^3 n^3)$ time. Such an approach also requires one to specify or learn kernel functions that bridge between the prior(s) and current Study, violating design goal 4.

Instead, our strategy is to build a stack of Gaussian Process regressors, where each regressor is associated with a study, and where each level is trained on the residuals relative to the regressor below it. Our model is that the studies were performed in a linear sequence, each study using the studies before it as priors.

The bottom of the stack contains a regressor built using data from the oldest study in the stack. The regressor above it is associated with the 2nd oldest study, and regresses on the *residual* of its




Figure 5: An illustration of our transfer learning scheme, showing how μ'_i is built from the residual labels w.r.t. μ_{i-1} (shown in dotted red lines).

objective relative to the predictions of the regressor below it. Similarly, the regressor associated with the i^{th} study is built using the data from that study, and regresses on the residual of the objective with respect to the predictions of the regressor below it.

More formally, we have a sequence of studies $\{S_i\}_{i=1}^k$ on unknown objective functions $\{f_i\}_{i=1}^k$, where the current study is S_k , and we build two sequences of regressors $\{R_i\}_{i=1}^k$ and $\{R'_i\}_{i=1}^k$ having posterior mean functions $\{\mu_i\}_{i=1}^k$ and $\{\mu'_i\}_{i=1}^k$ respectively, and posterior standard deviation functions $\{\sigma_i\}_{i=1}^k$ and $\{\sigma'_i\}_{i=1}^k$, respectively. Our final predictions will be μ_k and σ_k .

Let $D_i = \{(x_t^i, y_t^i)\}_t$ be the dataset for study S_i . Let R'_i be a regressor trained using data $\{(x_t^i, y_t^i - \mu_{i-1}(x_t^i))\}_t$ which computes μ'_i and σ'_i . Then we define as our posterior means at level i as $\mu_i(x) := \mu'_i(x) + \mu_{i-1}(x)$. We take our posterior standard deviations at level i , $\sigma_i(x)$, to be a weighted geometric mean of $\sigma'_i(x)$ and $\sigma_{i-1}(x)$, where the weights are a function of the amount of data (i.e., completed trials) in S_i and S_{i-1} . The exact weighting function depends on a constant $\alpha \approx 1$ sets the relative importance of old and new standard deviations.

This approach has nice properties when the prior regressors are densely supported (i.e. has many well-spaced data points), but the top-level regressor has relatively little training data: (1) fine structure in the priors carries through to μ_k , even if the top-level regressor gives a low-resolution model of the objective function residual; (2) since the estimate for σ'_k is inaccurate, averaging it with σ_{k-1} can lead to an improved estimate. Further, when the top-level regressor has dense support, $\beta \rightarrow 1$ and the $\sigma_k \rightarrow \sigma'_k$, as one might desire.

We provide details in the pseudocode in Algorithm 1, and illustrate the regressors in Figure 5.

Algorithm 1 is then used in the Batched Gaussian Process Bandits [8] algorithm. Algorithm 1 has the property that for a sufficiently dense sampling of the feasible region in the training data for the current study, the predictions converge to those of a regressor trained only on the current study data. This ensures a certain degree of robustness: badly chosen priors will eventually be overwhelmed (design goal 3).

Algorithm 1 Transfer Learning Regressor

```

# This is a higher order function that returns a regressor  $R(x_{\text{test}})$ ;
# then  $R(x_{\text{test}})$  can be evaluated to obtain  $(\mu, \sigma)$ 
function GETREGRESSOR( $D_{\text{training}}, i$ )
  If  $i < 0$ : Return function that returns (0,1) for all inputs
  # Recurse to get a Regressor  $(\mu_{i-1}(x), \sigma_{i-1}(x))$  trained on
  # the data for all levels of the stack below this one.
   $R_{\text{prior}} \leftarrow \text{GETREGRESSOR}(D_{\text{training}}, i - 1)$ 
  # Compute training residuals
   $D_{\text{residuals}} \leftarrow [(x, y - R_{\text{prior}}(x)[0]) \text{ for } (x, y) \in D_i]$ 
  # Train a Gaussian Process  $(\mu'_i(x), \sigma'_i(x))$  on the residuals.
   $GP_{\text{residuals}} = \text{TRAINGP}(D_{\text{residuals}})$ 
  function STACKEDREGRESSOR( $x_{\text{test}}$ )
     $\mu_{\text{prior}}, \sigma_{\text{prior}} \leftarrow R_{\text{prior}}(x_{\text{test}})$ 
     $\mu_{\text{top}}, \sigma_{\text{top}} \leftarrow GP_{\text{residuals}}(x_{\text{test}})$ 
     $\mu \leftarrow \mu_{\text{top}} + \mu_{\text{prior}}$ 
     $\beta \leftarrow \alpha |D_i| / (\alpha |D_i| + |D_{i-1}|)$ 
     $\sigma \leftarrow \sigma_{\text{top}}^\beta \sigma_{\text{prior}}^{1-\beta}$ 
    return  $\mu, \sigma$ 
  end function
  return STACKEDREGRESSOR
end function

```

In production settings, transfer learning is often particularly valuable when the number of trials per study is relatively small, but there are many such studies. For example, certain production machine learning systems may be very expensive to train, limiting the number of trials that can be run for hyperparameter tuning, yet are mission critical for a business and are thus worked on year after year. Over time, the total number of trials spanning several small hyperparameter tuning runs can be quite informative. Our transfer learning scheme is particularly well-suited to this case, as illustrated in section 4.3.

4 RESULTS

4.1 Performance Evaluation

To evaluate the performance of Google Vizier we require functions that can be used to benchmark the results. These are pre-selected, easily calculated functions with known optimal points that have proven challenging for black-box optimization algorithms. We can measure the success of an optimizer on a benchmark function f by its final *optimality gap*. That is, if x^* minimizes f , and \hat{x} is the best solution found by the optimizer, then $|f(\hat{x}) - f(x^*)|$ measures the success of that optimizer on that function. If, as is frequently the case, the optimizer has a stochastic component, we then calculate the *average optimality gap* by averaging over multiple runs of the optimizer on the same benchmark function.

Comparing between benchmarks is a more difficult given that the different benchmark functions have different ranges and difficulties. For example, a good black-box optimizer applied to the Rastrigin function might achieve an optimality gap of 160, while simple random sampling of the Beale function can quickly achieve an

optimality gap of 60 [10]. We normalize for this by taking the ratio of the optimality gap to the optimality gap of RANDOM SEARCH on the same function under the same conditions. Once normalized, we average over the benchmarks to get a single value representing an optimizer's performance.

The benchmarks selected were primarily taken from the Black-Box Optimization Benchmarking Workshop [10] (an academic competition for black-box optimizers), and include the Beale, Branin, Ellipsoidal, Rastrigin, Rosenbrock, Six Hump Camel, Sphere, and Styblinski benchmark functions.

4.2 Empirical Results

In Figures 6 we look at result quality for four optimization algorithms currently implemented in the Vizier framework: a multi-armed bandit technique using a Gaussian process regressor [29], the SMAC algorithm [19], the Covariance Matrix Adaption Evolution Strategy (CMA-ES) [16], and a probabilistic search method of our own. For a given dimension d , we generalized each benchmark function into a d dimensional space, ran each optimizer on each benchmark 100 times, and recorded the intermediate results (averaging these over the multiple runs). Figure 6 shows their improvement over RANDOM SEARCH; the horizontal axis represents the number of trials have been evaluated, while the vertical axis indicates each optimality gap as a fraction of the RANDOM SEARCH optimality gap at the same point. The 2×RANDOM SEARCH curve is the RANDOM SEARCH algorithm when it was allowed to sample two points for each point the other algorithms evaluated. While some authors have claimed that 2×RANDOM SEARCH is highly competitive with Bayesian Optimization methods [20], our data suggests this is only true when the dimensionality of the problem is sufficiently high (e.g., over 16).

4.3 Transfer Learning

We display the value of transfer learning in Figure 7 with a series of short studies; each study is just six trials long. Even so, one can see that transfer learning from one study to the next leads to steady progress towards the optimum, as the stack of regressors gradually builds up information about the shape of the objective function.

This experiment is conducted in a 10 dimensional space, using the 8 black-box functions described in section 4.1. We run 30 studies (180 trials) and each study uses transfer learning from all previous studies.

As one might hope, transfer learning causes the GP bandit algorithm to show a strong systematic decrease in the optimality gap from study to study, with its final average optimality gap 37% the size of RANDOM SEARCH's. As expected, RANDOM SEARCH shows no systematic improvement in its optimality gap from study to study.

Note that a systematic improvement in the optimality gap is a difficult task since each study gets a budget of only 6 trials whilst operating in a 10 dimensional space, and the GP regressor is optimizing 8 internal hyperparameters for each study. By any reasonable measure, a single study's data is insufficient for the regressor to learn much about the shape of the objective function.




Figure 6: Ratio of the average optimality gap of each optimizer to that of RANDOM SEARCH at a given number of samples. The 2XRANDOM SEARCH is a RANDOM SEARCH allowed to sample two points at every step (as opposed to a single point for the other algorithms).




Figure 7: Convergence of transfer learning in a 10 dimensional space. This shows a sequence of studies with progressive transfer learning for both GP Bandit (blue diamonds) and RANDOM SEARCH (red squares) optimizers. The X-axis shows the index of the study, i.e. the number of times that transfer learning has been applied; the Y-axis shows the log of the best mean optimality gap seen in the study (see Section 4.1). Each study contains six trials; for the GP Bandit-based optimizer the previous studies are used as priors for transfer learning. Note that the GP bandits shows a consistent improvement in optimality gap from study to study, thus demonstrating an effective transfer of knowledge from the earlier trials; RANDOM SEARCH does not do transfer learning.

4.4 Automated Stopping

4.4.1 Performance Curve Stopping Rule. In our experiments, we found that the use of the performance curve stopping rule resulted in achieving optimality gaps comparable to those achieved without the stopping rule, while using approximately 50% fewer CPU-hours when tuning hyperparameter for deep neural networks. Our result is in line with figures reported by other researchers, while using a

more flexible non-parametric model (e.g., Domhan et al. [9] report reductions in the 40% to 60% range on three ML hyperparameter tuning benchmarks).

4.4.2 Median Automated Stopping Rule. We evaluated the Median Stopping Rule for several hyperparameter search problems, including a state-of-the-art residual network architecture based on [17] for image classification on CIFAR10 with 16 tunable hyperparameters, and an LSTM architecture [33] for language modeling on the Penn TreeBank data set with 12 tunable hyperparameters. We observed that in all cases the stopping rule consistently achieved a factor two to three speedup over random search, while always finding the best performing Trial. Li et al. [20] argued that “2X random search”, i.e., random search at twice the speed, is competitive with several state-of-the-art black-box optimization methods on a broad range of benchmarks. The robustness of the stopping rule was also evaluated by running repeated simulations on a large set of completed random search trials under random permutation, which showed that the algorithm almost never decided to stop the ultimately-best-performing trial early.

5 USE CASES

Vizier is used for a number of different application domains.

5.1 Hyperparameter tuning and HyperTune

Vizier is used across Google to optimize hyperparameters of machine learning models, both for research and production models. Our implementation scales to service the entire hyperparameter tuning workload across Alphabet, which is extensive. As one (admittedly extreme) example, Collins et al. [6] used Vizier to perform hyperparameter tuning studies that collectively contained *millions* of trials for a research project investigating the capacity of different recurrent neural network architectures. In this context, a single trial involved training a distinct machine learning model using different hyperparameter values. That research project would not be possible without effective black-box optimization. For other research projects, automating the arduous and tedious task of hyperparameter tuning accelerates their progress.

Perhaps even more importantly, Vizier has made notable improvements to production models underlying many Google products, resulting in measurably better user experiences for over a billion people. External researchers and developers can achieve the same benefits using Google Cloud Machine Learning *HyperTune* subsystem, which benefits from our experience and technology.

5.2 Automated A/B testing

In addition to tuning hyperparameters, Vizier has a number of other uses. It is used for automated A/B testing of Google web properties, for example tuning user–interface parameters such as font and thumbnail sizes, color schema, and spacing, or traffic-serving parameters such as the relative importance of various signals in determining which items to show to a user. An example of the latter would be “how should the search results returned from Google Maps trade off search-relevance for distance from the user?”

5.3 Delicious Chocolate Chip Cookies

Vizier is also used to solve complex black–box optimization problems arising from physical design or logistical problems. Here we present an example that highlights some additional capabilities of the system: finding the most delicious chocolate chip cookie recipe from a parameterized space of recipes.

Parameters included baking soda, brown sugar, white sugar, butter, vanilla, egg, flour, chocolate, chip type, salt, cayenne, orange extract, baking time, and baking temperature. We provided recipes to contractors responsible for providing desserts for Google employees. The head chefs among the contractors were given discretion to alter parameters if (and only if) they strongly believed it to be necessary, but would carefully note what alterations were made. The cookies were baked, and distributed to the cafes for taste–testing. Cafe goers tasted the cookies and provided feedback via a survey. Survey results were aggregated and the results were sent back to Vizier. The “machine learning cookies” were provided about twice a week over several weeks.

The cookies improved significantly over time; later rounds were extremely well-rated and, in the authors’ opinions, delicious. However, we wish to highlight the following capabilities of Vizier the cookie design experiment exercised:

- **Infeasible trials:** In real applications, some trials may be *infeasible*, meaning they cannot be evaluated for reasons that are intrinsic to the parameter settings. Very high learning rates may cause training to diverge, leading to garbage models. In this example: very low levels of butter may make your cookie dough impossibly crumbly and incohesive.
- **Manual overrides of suggested trials:** Sometimes you cannot evaluate the suggested trial or else mistakenly evaluate a different trial than the one asked for. For example, when baking you might be running low on an ingredient and have to settle for less than the recommended amount.
- **Transfer learning:** Before starting to bake at large scale, we baked some recipes in a smaller scale run-through. This provided useful data that we could transfer learn from when

baking at scale. Conditions were not identical, however, resulting in some unexpected consequences. For example, the dough was allowed to sit longer in large-scale production which unexpectedly, and somewhat dramatically, increased the subjective spiciness of the cookies for trials that involved cayenne. Fortunately, our transfer learning scheme is relatively robust to such shifts.

Vizier supports marking trials as infeasible, in which case they do not receive an objective value. In the case of Bayesian Optimization, previous work either assigns them a particularly bad objective value, attempts to incorporate a probability of infeasibility into the acquisition function to penalize points that are likely to be infeasible [3], or tries to explicitly model the shape of the infeasible region [11, 12]. We take the first approach, which is simple and fairly effective for the applications we consider. Regarding manual overrides, Vizier’s stateless design makes it easy to support updating or deleting trials; we simply update the trial state on the database. For details on transfer learning, refer to section 3.3.

6 CONCLUSION

We have presented our design for Vizier, a scalable, state-of-the-art internal **service for black–box optimization** within Google, explained many of its design choices, and described its use cases and benefits. It has already proven to be a valuable platform for research and development, and we expect it will only grow more so as the area of black–box optimization grows in importance. Also, it designs excellent cookies, which is a very rare capability among computational systems.

7 ACKNOWLEDGMENTS

We gratefully acknowledge the contributions of the following: Jeremy Kubica, Jeff Dean, Eric Christiansen, Moritz Hardt, Katya Gonina, Kevin Jamieson, and Abdul Salem.

REFERENCES

- [1] Rémi Bardenet, Mátyás Brendel, Balázs Kégl, and Michele Sebag. 2013. Collaborative hyperparameter tuning. *ICML* 2 (2013), 199.
- [2] James S Bergstra, Rémi Bardenet, Yoshua Bengio, and Balázs Kégl. 2011. Algorithms for hyper-parameter optimization. In *Advances in Neural Information Processing Systems*. 2546–2554.
- [3] J Bernardo, MJ Bayarri, JO Berger, AP Dawid, D Heckerman, AFM Smith, and M West. 2011. Optimization under unknown constraints. *Bayesian Statistics* 9 9 (2011), 229.
- [4] Michael Bostock, Vadim Ogievetsky, and Jeffrey Heer. 2011. D³ data-driven documents. *IEEE transactions on visualization and computer graphics* 17, 12 (2011), 2301–2309.
- [5] Herman Chernoff. 1959. Sequential Design of Experiments. *Ann. Math. Statist.* 30, 3 (09 1959), 755–770. <https://doi.org/10.1214/aoms/1177706205>
- [6] Jasmine Collins, Jascha Sohl-Dickstein, and David Sussillo. 2017. Capacity and Trainability in Recurrent Neural Networks. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- [7] Andrew R Conn, Katya Scheinberg, and Luis N Vicente. 2009. *Introduction to derivative-free optimization*. SIAM.
- [8] Thomas Desautels, Andreas Krause, and Joel W Burdick. 2014. Parallelizing exploration-exploitation tradeoffs in Gaussian process bandit optimization. *Journal of Machine Learning Research* 15, 1 (2014), 3873–3923.
- [9] Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. 2015. Speeding Up Automatic Hyperparameter Optimization of Deep Neural Networks by Extrapolation of Learning Curves. In *IJCAI*. 3460–3468.

- [10] Steffen Finck, Nikolaus Hansen, Raymond Rost, and Anne Auger. 2009. Real-Parameter Black-Box Optimization Benchmarking 2009: Presentation of the Noiseless Functions. <http://coco.gforge.inria.fr/lib/exe/fetch.php?media=download3.6:bbobdocfunctions.pdf>. (2009). [Online].
- [11] Jacob R Gardner, Matt J Kusner, Zhixiang Eddie Xu, Kilian Q Weinberger, and John P Cunningham. 2014. Bayesian Optimization with Inequality Constraints. In *ICML*. 937–945.
- [12] Michael A Gelbart, Jasper Snoek, and Ryan P Adams. 2014. Bayesian optimization with unknown constraints. In *Proceedings of the Thirtieth Conference on Uncertainty in Artificial Intelligence*. AUAI Press, 250–259.
- [13] Josep Ginebra and Murray K. Clayton. 1995. Response Surface Bandits. *Journal of the Royal Statistical Society. Series B (Methodological)* 57, 4 (1995), 771–784. <http://www.jstor.org/stable/2345943>
- [14] Google. 2017. Polymer: Build modern apps using web components. <https://github.com/Polymer/polymer>. (2017). [Online].
- [15] Google. 2017. Protocol Buffers: Google's data interchange format. <https://github.com/google/protobuf>. (2017). [Online].
- [16] Nikolaus Hansen and Andreas Ostermeier. 2001. Completely derandomized self-adaptation in evolution strategies. *Evolutionary computation* 9, 2 (2001), 159–195.
- [17] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 770–778.
- [18] Julian Heinrich and Daniel Weiskopf. 2013. State of the Art of Parallel Coordinates. In *Eurographics (STARs)*. 95–116.
- [19] Frank Hutter, Holger H Hoos, and Kevin Leyton-Brown. 2011. Sequential model-based optimization for general algorithm configuration. In *International Conference on Learning and Intelligent Optimization*. Springer, 507–523.
- [20] Lisha Li, Kevin G. Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. 2016. Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization. *CoRR* abs/1603.06560 (2016). <http://arxiv.org/abs/1603.06560>
- [21] J Mockus, V Tiesis, and A Zilinskas. 1978. *The Application of Bayesian Methods for Seeking the Extremum*. Vol. 2. Elsevier. 117–128 pages.
- [22] John A Nelder and Roger Mead. 1965. A simplex method for function minimization. *The computer journal* 7, 4 (1965), 308–313.
- [23] Carl Edward Rasmussen and Christopher K. I. Williams. 2005. *Gaussian Processes for Machine Learning (Adaptive Computation and Machine Learning)*. The MIT Press.
- [24] Luis Miguel Rios and Nikolaos V Sahinidis. 2013. Derivative-free optimization: a review of algorithms and comparison of software implementations. *Journal of Global Optimization* 56, 3 (2013), 1247–1293.
- [25] Bobak Shahriari, Kevin Swersky, Ziyu Wang, Ryan P Adams, and Nando de Freitas. 2016. Taking the human out of the loop: A review of bayesian optimization. *Proc. IEEE* 104, 1 (2016), 148–175.
- [26] Jasper Snoek, Hugo Larochelle, and Ryan P Adams. 2012. Practical bayesian optimization of machine learning algorithms. In *Advances in neural information processing systems*. 2951–2959.
- [27] Jasper Snoek, Oren Rippel, Kevin Swersky, Ryan Kiros, Nadathur Satish, Narayanan Sundaram, Md. Mostofa Ali Patwary, Prabhat, and Ryan P. Adams. 2015. Scalable Bayesian Optimization Using Deep Neural Networks. In *Proceedings of the 32nd International Conference on Machine Learning, ICML 2015, Lille, France, 6-11 July 2015 (JMLR Workshop and Conference Proceedings)*, Francis R. Bach and David M. Blei (Eds.), Vol. 37. JMLR.org, 2171–2180. <http://jmlr.org/proceedings/papers/v37/snoek15.html>
- [28] Jost Tobias Springenberg, Aaron Klein, Stefan Falkner, and Frank Hutter. 2016. Bayesian Optimization with Robust Bayesian Neural Networks. In *Advances in Neural Information Processing Systems 29*, D. D. Lee, M. Sugiyama, U. V. Luxburg, I. Guyon, and R. Garnett (Eds.). Curran Associates, Inc., 4134–4142. <http://papers.nips.cc/paper/6117-bayesian-optimization-with-robust-bayesian-neural-networks.pdf>
- [29] Niranjan Srinivas, Andreas Krause, Sham Kakade, and Matthias Seeger. 2010. Gaussian Process Optimization in the Bandit Setting: No Regret and Experimental Design. *ICML* (2010).
- [30] Kevin Swersky, Jasper Snoek, and Ryan Prescott Adams. 2014. Freeze-thaw Bayesian optimization. *arXiv preprint arXiv:1406.3896* (2014).
- [31] Andrew Gordon Wilson, Zhiting Hu, Ruslan Salakhutdinov, and Eric P Xing. 2016. Deep kernel learning. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*. 370–378.
- [32] Dani Yogatama and Gideon Mann. 2014. Efficient Transfer Learning Method for Automatic Hyperparameter Tuning. *JMLR: W&CP* 33 (2014), 1077–1085.
- [33] Wojciech Zaremba, Ilya Sutskever, and Oriol Vinyals. 2014. Recurrent neural network regularization. *arXiv preprint arXiv:1409.2329* (2014).